

Siddhant Waghjale

swaghjal@andrew.cmu.edu | (412) 390-6437 | LinkedIn: [siddhantwaghjale](#)

EDUCATION

Carnegie Mellon University, School of Computer Science

Pittsburgh, PA

Master of Science in Intelligent Information Systems (Machine Learning and Natural Language Processing), GPA: 4.06/4.0

December 2024

- Coursework: Introduction to Machine Learning, Advanced Natural Language Processing, Question Answering, Multimodal Machine Learning, Neural Code-Generation, Visual Learning & Representations

National Institute of Technology, Karnataka

Karnataka, India

Bachelor of Technology in Information Technology, GPA: 8.17/10

June 2020

EXPERIENCE

Tesla

Palo Alto, USA

Autopilot | Software Intern

May 2024 - August 2024

- Developed an AI-Oncall ChatBot that resolved 50% of user queries without human intervention by implementing advanced RAG techniques, real-time vector database updates, and multi-channel integration for 300 software engineers.
- Created a Failure Summary feature that reduced build issue resolution time by 80% through developing an LLM agent-based summarizer, integrating it into the Builds UI.
- Enhanced SDK performance by reducing download and extraction time 60x through improving caching logic and optimizing the extraction process.

Kaleyra

Bangalore, India

Chatbot Team | Senior Associate Software Engineer (Team Lead)

September 2020 - June 2023

- Led R&D team in developing a scalable Chatbot with advanced NLU using Transformers, achieving 350 TPS and serving 10M customers daily while reducing AWS costs by \$823/month
- Engineered Email and SMS Spam Detection Engines by fine-tuning BERT and optimizing with ONNX for a 20x boost in classification latency and throughput.
- Devised an automatic authentication system with document and face verification (AWS ISV Innovation Cup finalist).

Deloitte Consulting USI

Bangalore, India

Oracle Cloud Solutions | Software Intern

May 2019 - July 2019

- Developed expertise in real-time business processes, BI reporting, and data modeling by creating data models, formulating PL/SQL queries, and achieving key project milestones.

PUBLICATIONS

ECCO: Can We Improve Model-Generated Code Efficiency Without Sacrificing Functional Correctness?

(Under Review)

Siddhant Waghjale, Vishruth Veerendranath, Zora Zhiruo Wang, Daniel Fried

ACADEMIC PROJECTS

Two-Stage Multimodal Architecture for Visual Abductive Reasoning

Multimodal Visual Commonsense Reasoning

Jan 2024 - April 2024

- Developed a two-stage training architecture for abductive visual reasoning, enhancing image captioning models to infer higher-level intents and social dynamics.
- Fine-tuned language and vision models on targeted multimodal data, achieving significant gains on the Sherlock dataset.

Improving Performance for the Multi-Table Question Answering Task

Multi-Table Question Answering

Jan 2024 - April 2024

- Implemented two approaches for multi-table QA using code-based models and chain-of-table reasoning
- Improved evaluation metrics and conducted extensive quantitative and qualitative analysis to enhance performance.

Detecting LLM Generated Text in Multi-generator, Multi-domain, and Multi-lingual Black-Box Setting

SemEval-2024 Task 8: Multigenerator, Multidomain, and Multilingual Black-Box Machine-Generated Text Detection

August 2023 - December 2023

- Optimized Mistral 7B and Llama2 7B on the M4 dataset for monolingual and multilingual text detection, achieving peak accuracies of 92% and 77%, respectively
- Performed paraphrase tests, showing that the Logistic Regression model generalized better than fine-tuned models.

SKILLS

Languages: Python, Golang, C++

Libraries & Frameworks: PyTorch, TensorFlow, Keras, Langchain, Scikit-Learn, Huggingface, Spacy, OpenCV, Flask, FastAPI, ONNX, CUDA

Databases & Message-Brokers: MongoDB, Redis, SQL, Object Storage, ChromaDB, Snowflake, PostgreSQL, RabbitMQ